



清华大学  
Tsinghua University

Advanced Computer Vision  
THU×SENSETIME – 80231202



Chapter 2 - Section 8

# Image Segmentation

Dr. Li Hongyang

Friday, April 15, 2022

Partial credit by : Wang Cheng

# Outline

- 
- Part 1**      **Recap: classification loss and detection pipeline**
- 
- Part 2**      **3D Detection and BEV Perception**
- 
- Part 3**      **Image segmentation**

- Classification Loss

- Cross entropy or MSE?

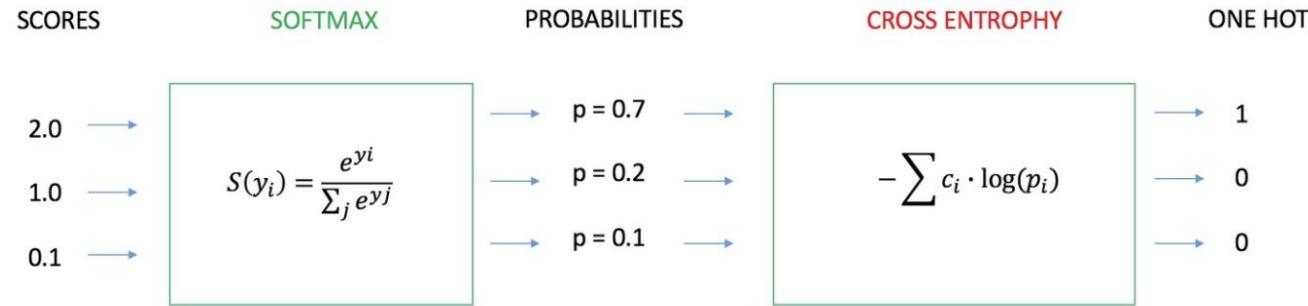
For MSE loss, we have

$$L_{mse,i} = \frac{1}{2}(y_i - p_i)^2$$

$p_i = \sigma(Wx + b)$  函数  $\sigma(\cdot)$  是 Sigmoid 函数

The grad for MSE loss is:  $\partial L_{mse,i} / \partial W = (y_i - p_i)p_i(1 - p_i)x$

- $Wx+b$  could be very large/small in the first few iterations; since parameters are randomly initialized.
- Gradient vanishing!

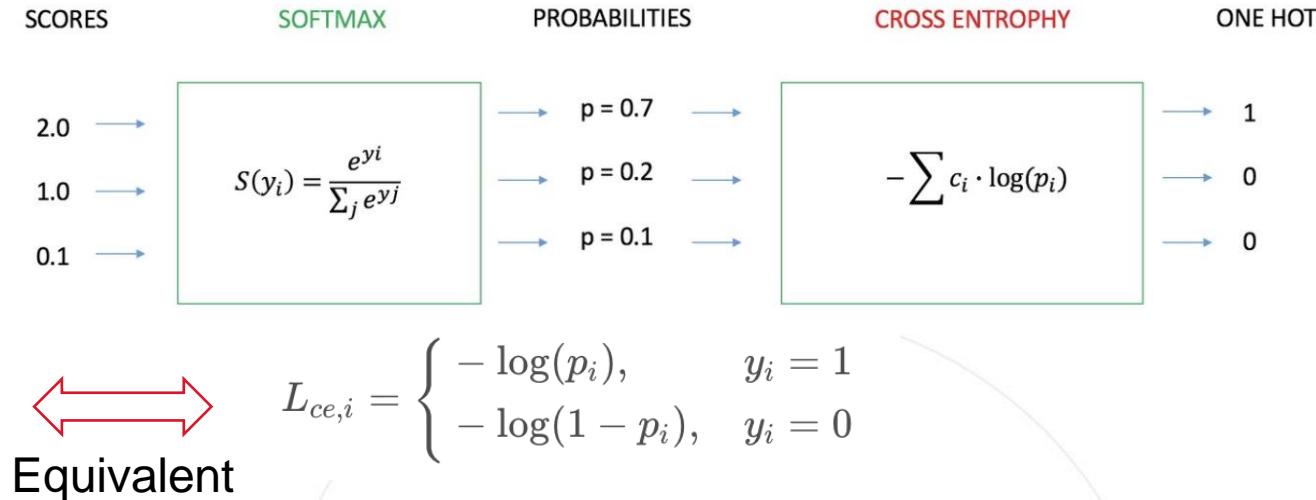


- Classification Loss

- Cross entropy or MSE?

For Cross entropy loss, we have

$$L_{ce,i} = -[y_i \log(p_i) + (1 - y_i) \log(1 - p_i)]$$



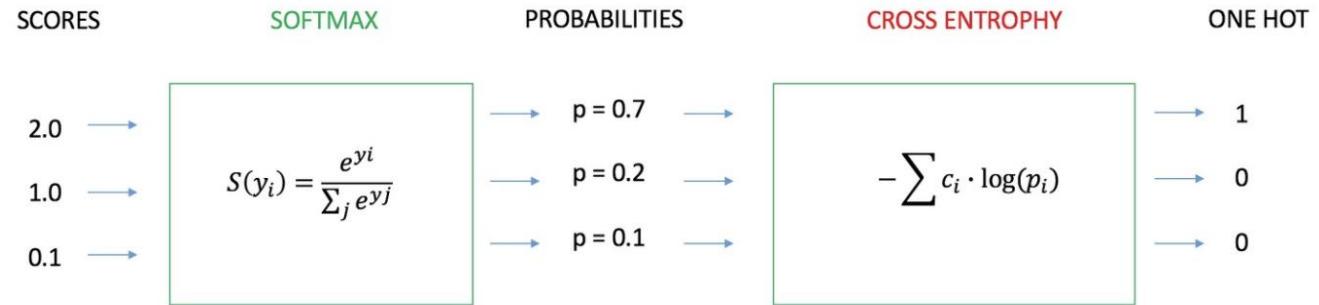
$$L_{ce,i} = \begin{cases} -\log(p_i), & y_i = 1 \\ -\log(1 - p_i), & y_i = 0 \end{cases}$$

The grad for CE loss is:  $\partial L_{ce,i} / \partial W = -(y_i - p_i)x$

- Gradient vanishing! resolved

- Classification Loss

- Cross entropy or MSE?



Multi-classification CE loss:

$$p_{i,k} = \exp(q_{i,k}) / \sum_j^M \exp(q_{i,j})$$

$$L = \sum_i^N L_{ce,i} = - \sum_i^N \sum_j^M y_{i,j} \log(p_{i,j})$$

**Further extension:**

From cross entropy loss to Focal loss and Circle loss  
What's the relationship?

- Background

- Consistency Regularization

$$\|p_{\text{model}}(y \mid \text{Augment}(x); \theta) - p_{\text{model}}(y \mid x; \theta)\|_2^2$$

模型在无标签数据 **增广前** 和 **增广后** 的预测应该一致

- Entropy Minimization

- Minimizes the entropy of  $p_{\text{model}}(y|x; \theta)$

要求分类器对无标签样本输出熵较少的结果

- Traditional Regularization

加入一些常量干扰，使得模型难以记住训练样本，来达到泛化要求

- **Mixup**: Beyond Empirical Risk Minimization (经验风险最小化；训练误差越小越好)
  - A simple and data-agnostic data augmentation method

$$\tilde{x} = \lambda x_i + (1 - \lambda)x_j,$$

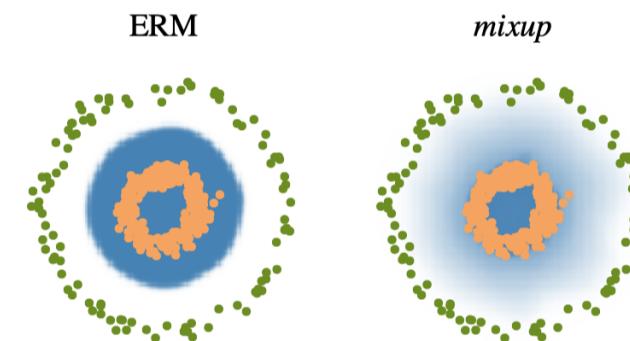
where  $x_i, x_j$  are raw input vectors

$$\tilde{y} = \lambda y_i + (1 - \lambda)y_j,$$

where  $y_i, y_j$  are one-hot label encodings

```
# y1, y2 should be one-hot vectors
for (x1, y1), (x2, y2) in zip(loader1, loader2):
    lam = numpy.random.beta(alpha, alpha)
    x = Variable(lam * x1 + (1. - lam) * x2)
    y = Variable(lam * y1 + (1. - lam) * y2)
    optimizer.zero_grad()
    loss(net(x), y).backward()
    optimizer.step()
```

(a) One epoch of *mixup* training in PyTorch.



(b) Effect of *mixup* ( $\alpha = 1$ ) on a toy problem. Green: Class 0. Orange: Class 1. Blue shading indicates  $p(y = 1|x)$ .

- mixup: Beyond Empirical Risk Minimization
  - Why use beta distribution?
  - Similar to label smooth?
  - Why mixup works? (Generalization gap between training data and real data)

<https://www.zhihu.com/question/67472285>

## 如何评价mixup: BEYOND EMPIRICAL RISK MINIMIZATION?

这篇paper是ICLR2018的投稿，直接对raw data和label interpolation，在很多数据集上取得了SoTA。arxiv.org/pdf/1710.0941...

关注问题 写回答 邀请回答 好问题 4 添加评论 分享 收藏

20个回答

默认排序

张宏毅  
MIT 机器学习 / 认知科学  
456人赞同了该回答



谢流远



深度学习 (Deep Learning) 话题下的优秀答主

77人赞同了该回答

Mixup超好用的，轻松提高一个点，参见我们的paper：

[arxiv.org/abs/1812.0118...](https://arxiv.org/abs/1812.0118...)

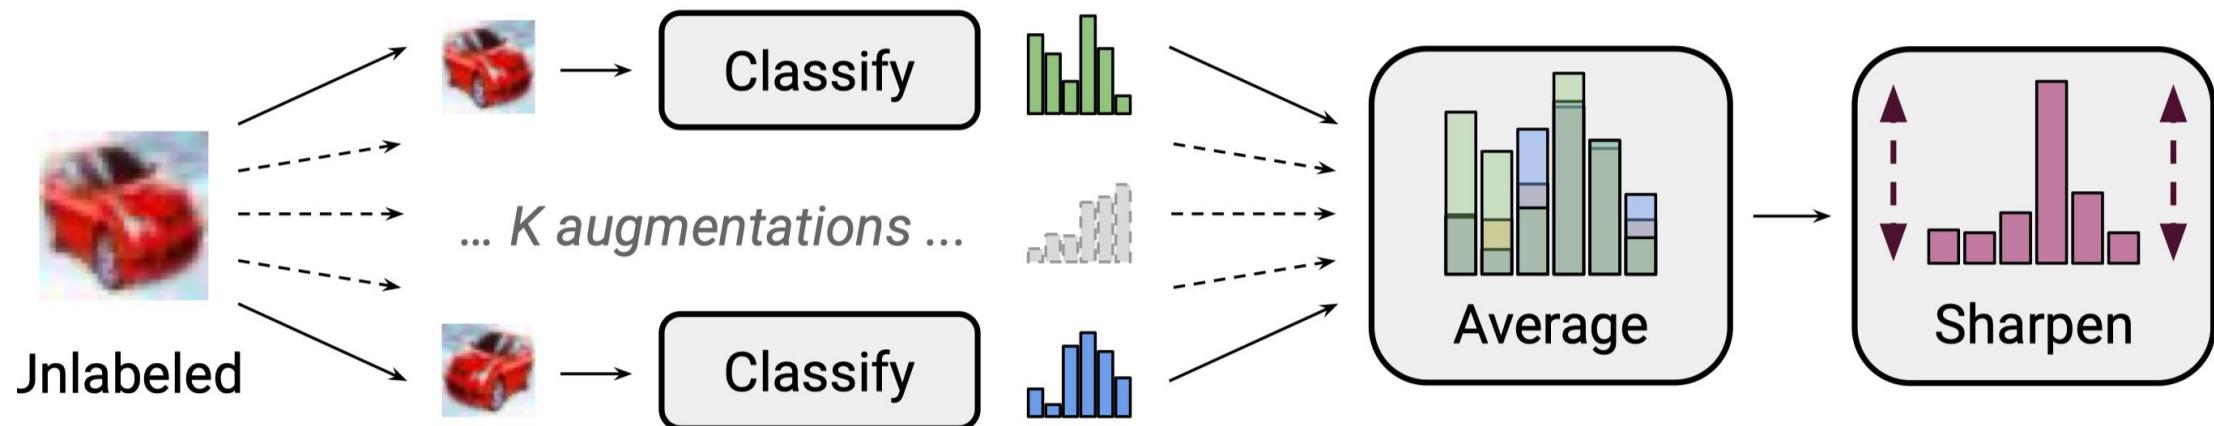
编辑于 2018-12-07

- mixup: Beyond Empirical Risk Minimization

## Further reading

- MixMatch
- ReMixMatch
- FixMatch

- MixMatch: A Holistic Approach to Semi-Supervised Learning
  - Stochastic data augmentation is applied to an unlabeled image K times
  - The average of these K predictions is “sharpened” by adjusting the distribution’s temperature



Berthelot D, Carlini N, Goodfellow I, Papernot N, Oliver A, Raffel C. Mixmatch: A holistic approach to semi-supervised learning. arXiv preprint arXiv:1905.02249. 2019 May 6.

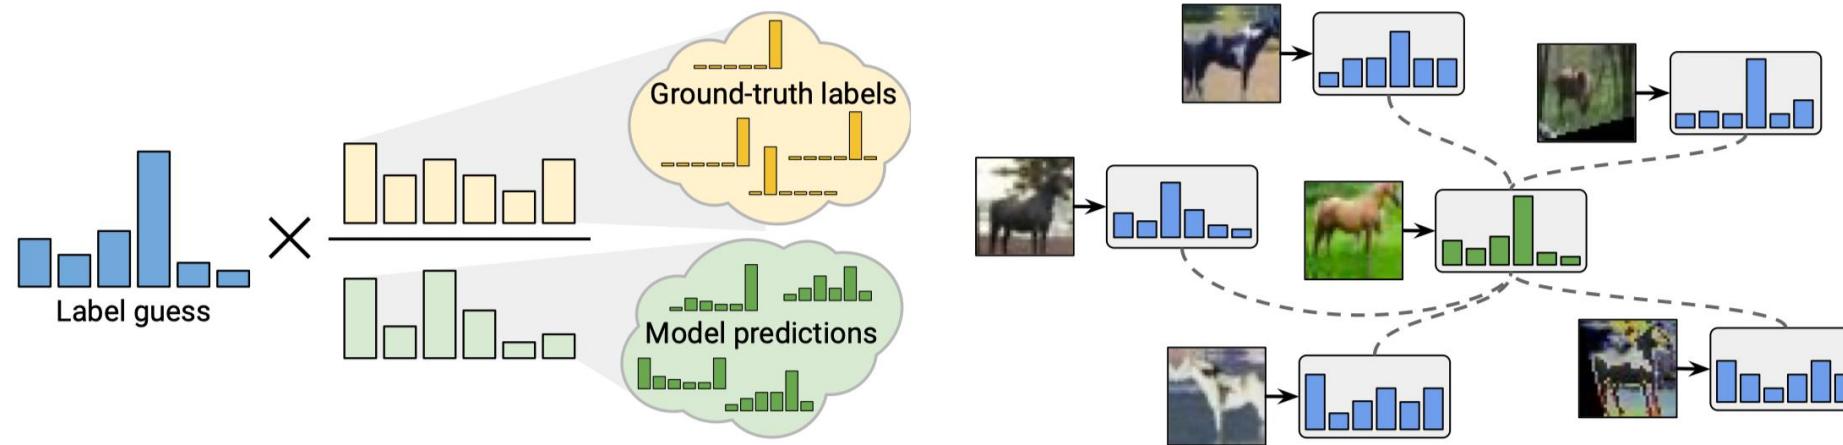
- MixMatch: A Holistic Approach to Semi-Supervised Learning

```

1: Input: Batch of labeled examples and their one-hot labels  $\mathcal{X} = ((x_b, p_b); b \in (1, \dots, B))$ , batch of unlabeled examples  $\mathcal{U} = (u_b; b \in (1, \dots, B))$ , sharpening temperature  $T$ , number of augmentations  $K$ , Beta distribution parameter  $\alpha$  for MixUp.
2: for  $b = 1$  to  $B$  do
3:    $\hat{x}_b = \text{Augment}(x_b)$  // Apply data augmentation to  $x_b$ 
4:   for  $k = 1$  to  $K$  do
5:      $\hat{u}_{b,k} = \text{Augment}(u_b)$  // Apply  $k^{th}$  round of data augmentation to  $u_b$ 
6:   end for
7:    $\bar{q}_b = \frac{1}{K} \sum_k p_{\text{model}}(y | \hat{u}_{b,k}; \theta)$  // Compute average predictions across all augmentations of  $u_b$ 
8:    $q_b = \text{Sharpen}(\bar{q}_b, T)$  // Apply temperature sharpening to the average prediction (see eq. (7))
9: end for
10:   $\hat{\mathcal{X}} = ((\hat{x}_b, p_b); b \in (1, \dots, B))$  // Augmented labeled examples and their labels
11:   $\hat{\mathcal{U}} = ((\hat{u}_{b,k}, q_b); b \in (1, \dots, B), k \in (1, \dots, K))$  // Augmented unlabeled examples, guessed labels
12:   $\mathcal{W} = \text{Shuffle}(\text{Concat}(\hat{\mathcal{X}}, \hat{\mathcal{U}}))$  // Combine and shuffle labeled and unlabeled data
13:   $\mathcal{X}' = (\text{MixUp}(\hat{\mathcal{X}}_i, \mathcal{W}_i); i \in (1, \dots, |\hat{\mathcal{X}}|))$  // Apply MixUp to labeled data and entries from  $\mathcal{W}$ 
14:   $\mathcal{U}' = (\text{MixUp}(\hat{\mathcal{U}}_i, \mathcal{W}_{i+|\hat{\mathcal{X}}|}); i \in (1, \dots, |\hat{\mathcal{U}}|))$  // Apply MixUp to unlabeled data and the rest of  $\mathcal{W}$ 
15: return  $\mathcal{X}', \mathcal{U}'$ 

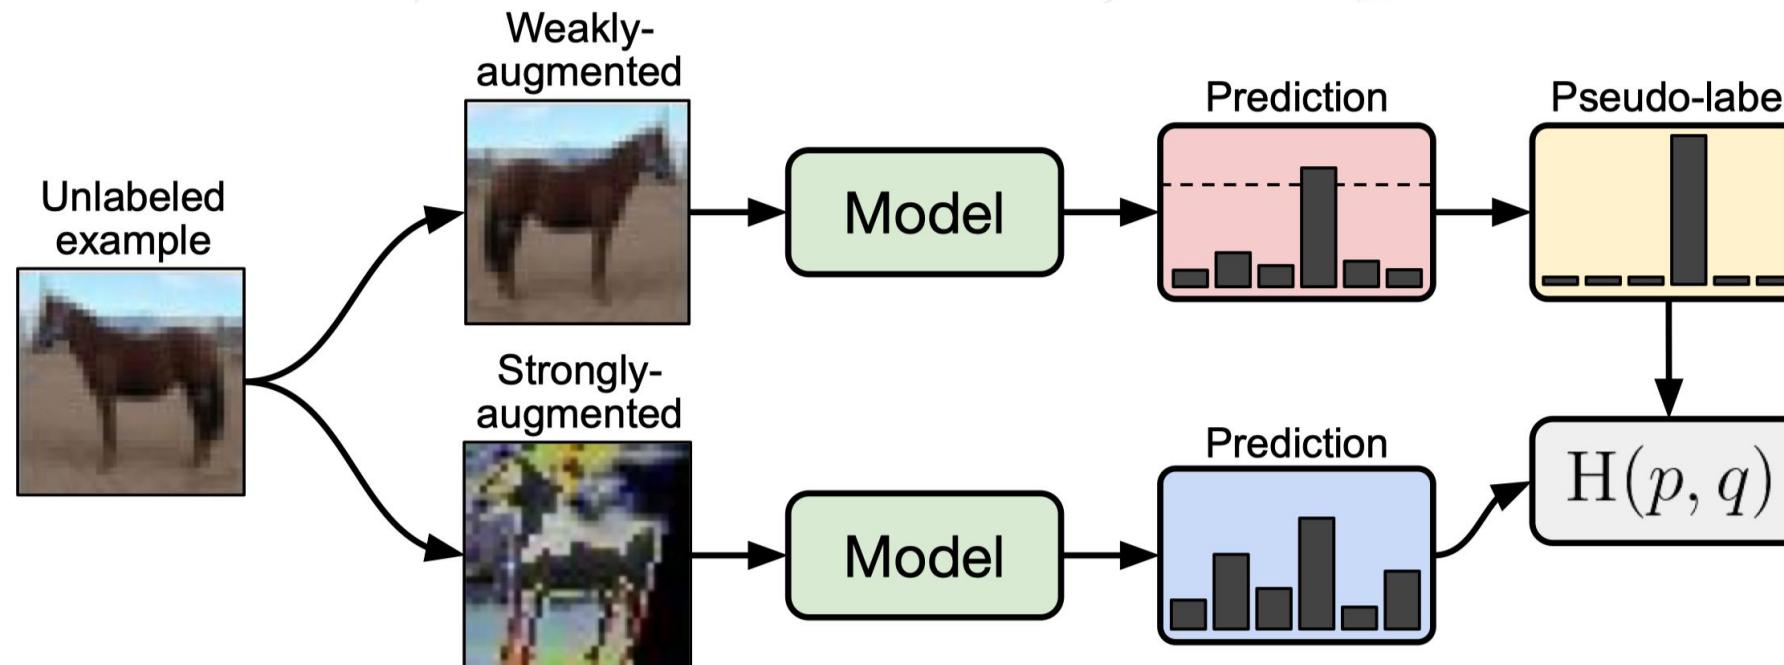
```

- ReMixMatch: Semi-Supervised Learning with Distribution Matching and Augmentation Anchoring
  - Improved version of MixMatch
  - Distribution Alignment (left) and Augmentation Anchor (right)



Berthelot D, Carlini N, et al. Remixmatch: Semi-supervised learning with distribution alignment and augmentation anchoring. arXiv preprint arXiv:1911.09785. 2019 Nov 21.

- FixMatch: Simplifying Semi-Supervised Learning with Consistency and Confidence
  - Combination of Consistency regularization and pseudo-labeling.



Sohn K, Berthelot D, et al. Fixmatch: Simplifying semi-supervised learning with consistency and confidence. arXiv preprint arXiv:2001.07685. 2020 Jan 21.

- FixMatch: Simplifying Semi-Supervised Learning with Consistency and Confidence
  - FixMatch consists of two loss terms: a supervised loss  $\ell_s$  and an unsupervised loss  $\ell_u$
  - $\ell_s$  is the standard cross-entropy loss

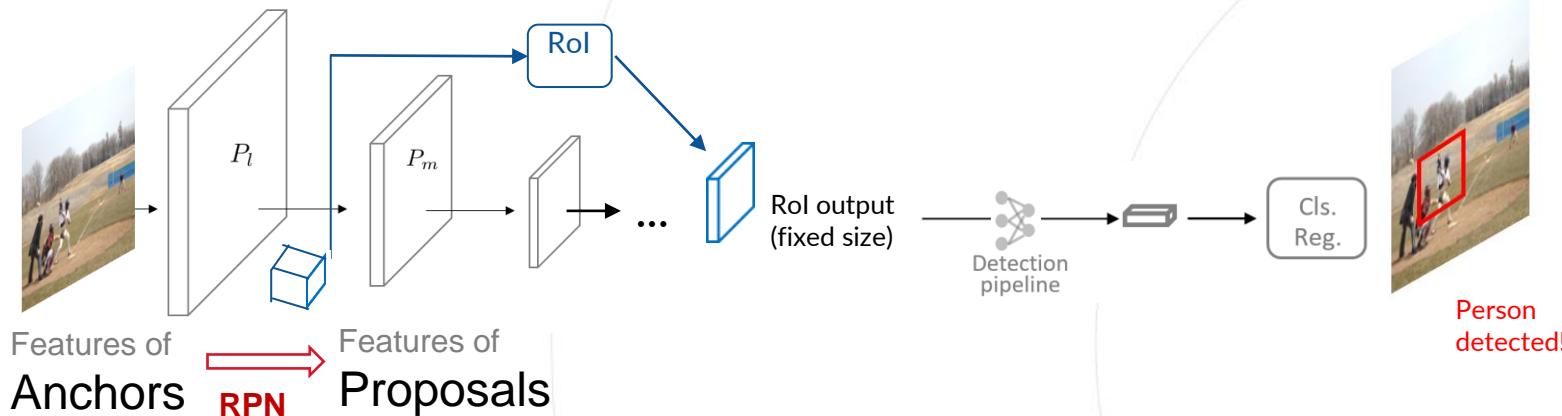
$$\ell_s = \frac{1}{B} \sum_{b=1}^B H(p_b, p_m(y | \alpha(x_b)))$$

- Convert the prediction on the weakly-augmented image to a one-hot pseudo-label
- $\ell_u$  is the cross-entropy loss against the model's output for the strongly-augmented image

$$\ell_u = \frac{1}{\mu B} \sum_{b=1}^{\mu B} \mathbb{1}(\max(q_b) \geq \tau) H(\hat{q}_b, p_m(y | \mathcal{A}(u_b)))$$

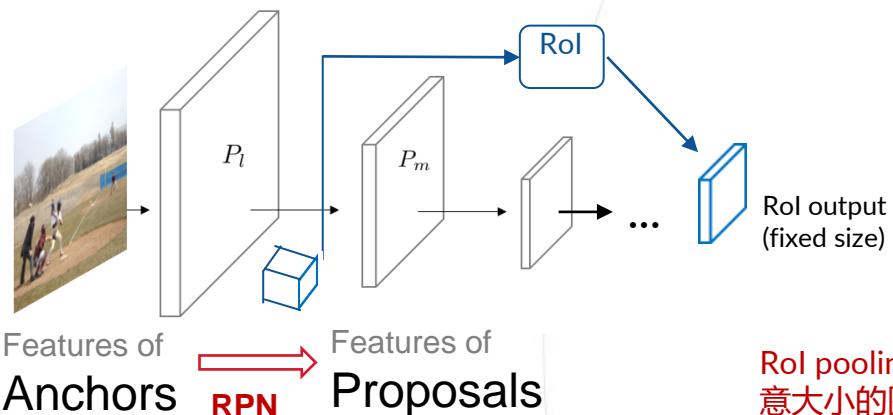
- **Two stage** vs One-stage pipeline

- Anchor placement: large anchors should be put in the low-level layers or high-level layers?



- **Two stage** vs One-stage pipeline

- Loss of RoI pooling



RoI pooling 的出现，让我们能用任意大小的图像作为输入，总能产生固定大小的输出。

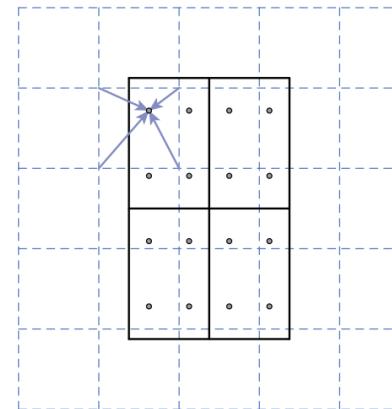
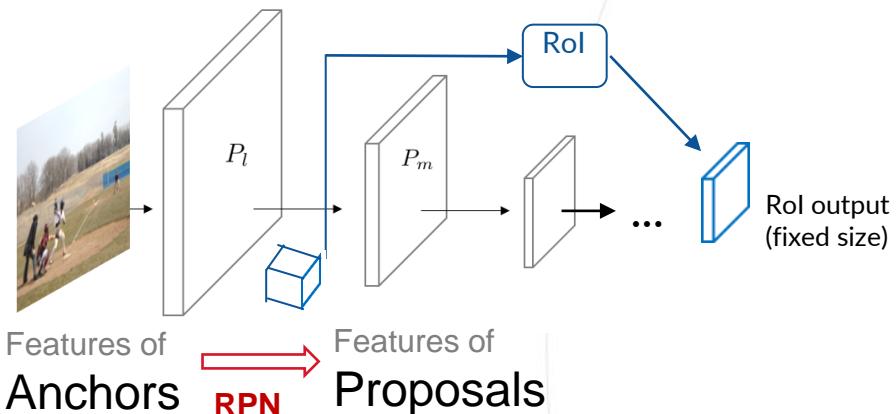
$$\frac{\partial L}{\partial x_i} = \sum_r \sum_j [i = i^*(r, j)] \frac{\partial L}{\partial y_{rj}}. \quad (4)$$

In words, for each mini-batch RoI  $r$  and for each pooling output unit  $y_{rj}$ , the partial derivative  $\partial L / \partial y_{rj}$  is accumulated if  $i$  is the argmax selected for  $y_{rj}$  by max pooling. In back-propagation, the partial derivatives  $\partial L / \partial y_{rj}$  are already computed by the backwards function of the layer on top of the RoI pooling layer.

RoI layer的BP计算。  
详见Fast RCNN paper.

- **Two stage** vs One-stage pipeline

- Loss of RoI pooling



**Figure 3. RoIAlign:** The dashed grid represents a feature map, the solid lines an ROI (with  $2 \times 2$  bins in this example), and the dots the 4 sampling points in each bin. RoIAlign computes the value of each sampling point by bilinear interpolation from the nearby grid points on the feature map. No quantization is performed on any coordinates involved in the ROI, its bins, or the sampling points.

- Two stage vs **One-stage pipeline**

- No RPN, No RoI-pooling

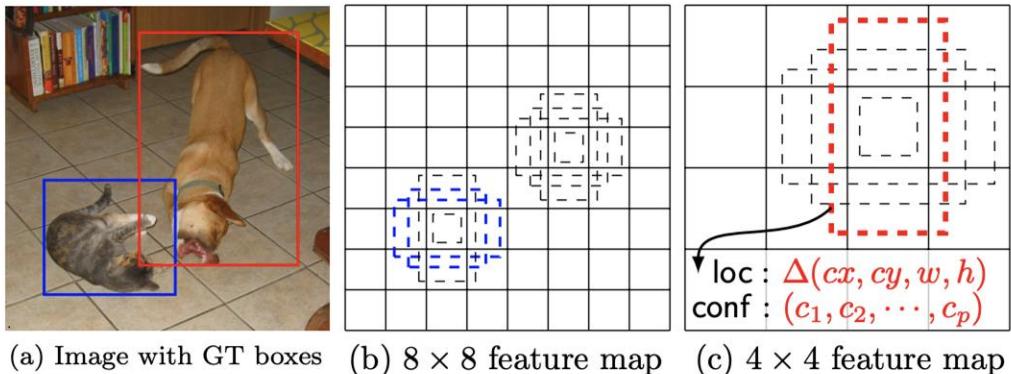


Fig. 1: **SSD framework.** (a) SSD only needs an input image and ground truth boxes for each object during training. In a convolutional fashion, we evaluate a small set (e.g. 4) of default boxes of different aspect ratios at each location in several feature maps with different scales (e.g.  $8 \times 8$  and  $4 \times 4$  in (b) and (c)). For each default box, we predict both the shape offsets and the confidences for all object categories ( $(c_1, c_2, \dots, c_p)$ ). At training time, we first match these default boxes to the ground truth boxes. For example, we have matched two default boxes with the cat and one with the dog, which are treated as positives and the rest as negatives. The model loss is a weighted sum between localization loss (e.g. Smooth L1 [6]) and confidence loss (e.g. Softmax).

What's the output feature map in this (c) example?

Answer:  $4 \times 4 \times (4 * 4 + p)$

Within **each grid cell**:

- Regress from each for the **B base boxes** (aka anchors) to a final box with **(dx, dy, dh, dw)**
- Predict scores for each of **p classes**

# Recap: Image Detection

- Detection loss

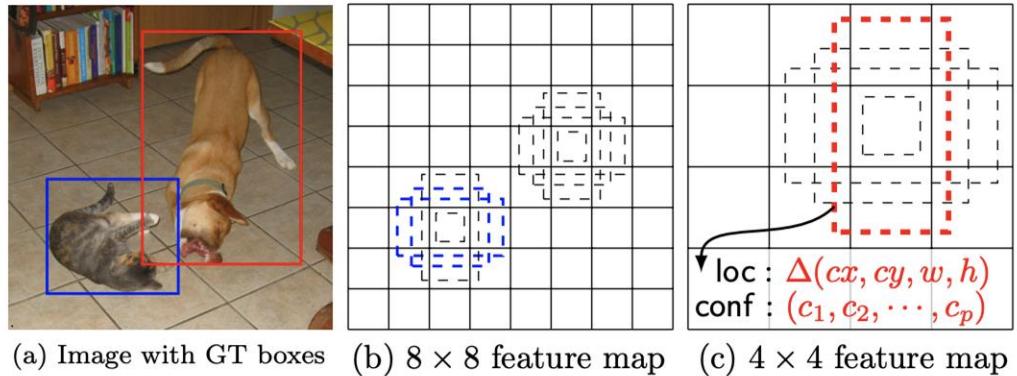


Fig. 1: **SSD framework.** (a) SSD only needs an input image and ground truth boxes for each object during training. In a convolutional fashion, we evaluate a small set (e.g. 4) of default boxes of different aspect ratios at each location in several feature maps with different scales (e.g.  $8 \times 8$  and  $4 \times 4$  in (b) and (c)). For each default box, we predict both the shape offsets and the confidences for all object categories ( $(c_1, c_2, \dots, c_p)$ ). At training time, we first match these default boxes to the ground truth boxes. For example, we have matched two default boxes with the cat and one with the dog, which are treated as positives and the rest as negatives. The model loss is a weighted sum between localization loss (e.g. Smooth L1 [6]) and confidence loss (e.g. Softmax).

<https://arxiv.org/pdf/1506.01497.pdf>

$$L(\{p_i\}, \{t_i\}) = \frac{1}{N_{cls}} \sum_i L_{cls}(p_i, p_i^*) + \lambda \frac{1}{N_{reg}} \sum_i p_i^* L_{reg}(t_i, t_i^*).$$

$p^*$  and  $t^*$  are the ground truth for classification and localization/regression

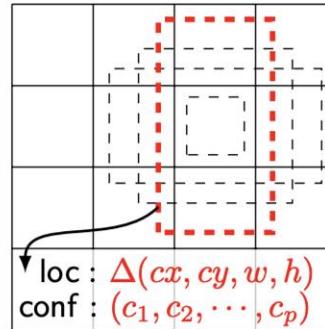
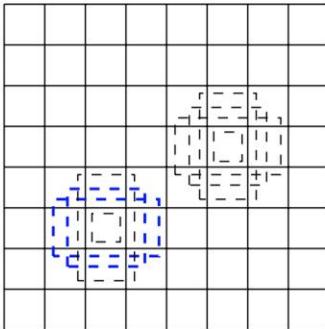
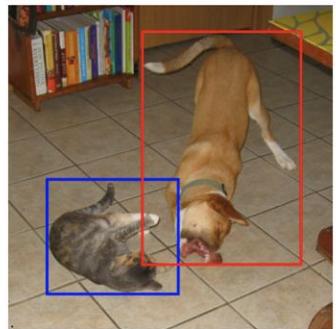
- Note that regression loss is only for positive samples.

$$L_{reg}(t_i, t_i^*) = R(t_i - t_i^*)$$

$R$  is the smoothed L1 loss

# Recap: Image Detection

- Detection loss



A bounding-box regression from  
**an anchor box** to a nearby  
ground-truth box.

$$\begin{aligned} t_x &= (x - x_a)/w_a, & t_y &= (y - y_a)/h_a, \\ t_w &= \log(w/w_a), & t_h &= \log(h/h_a), \\ t_x^* &= (x^* - x_a)/w_a, & t_y^* &= (y^* - y_a)/h_a, \\ t_w^* &= \log(w^*/w_a), & t_h^* &= \log(h^*/h_a), \end{aligned}$$

<https://arxiv.org/pdf/1506.01497.pdf>

$$L(\{p_i\}, \{t_i\}) = \frac{1}{N_{cls}} \sum_i L_{cls}(p_i, p_i^*)$$

$$+ \lambda \frac{1}{N_{reg}} \sum_i p_i^* L_{reg}(t_i, t_i^*).$$

$p^*$  and  $t^*$  are the ground truth for classification and localization/regression

- Note that regression loss is only for positive samples.

$$L_{reg}(t_i, t_i^*) = R(t_i - t_i^*)$$

$R$  is the smoothed L1 loss

$$\text{smooth}_{L_1}(x) = \begin{cases} 0.5x^2 & \text{if } |x| < 1 \\ |x| - 0.5 & \text{otherwise,} \end{cases}$$

# Recap: Image Detection

- YOLO vs SSD

SSD:

- Smaller input size
- Faster FPS

SSD:

- More templates/anchors from various depth in the network
- Higher mAP

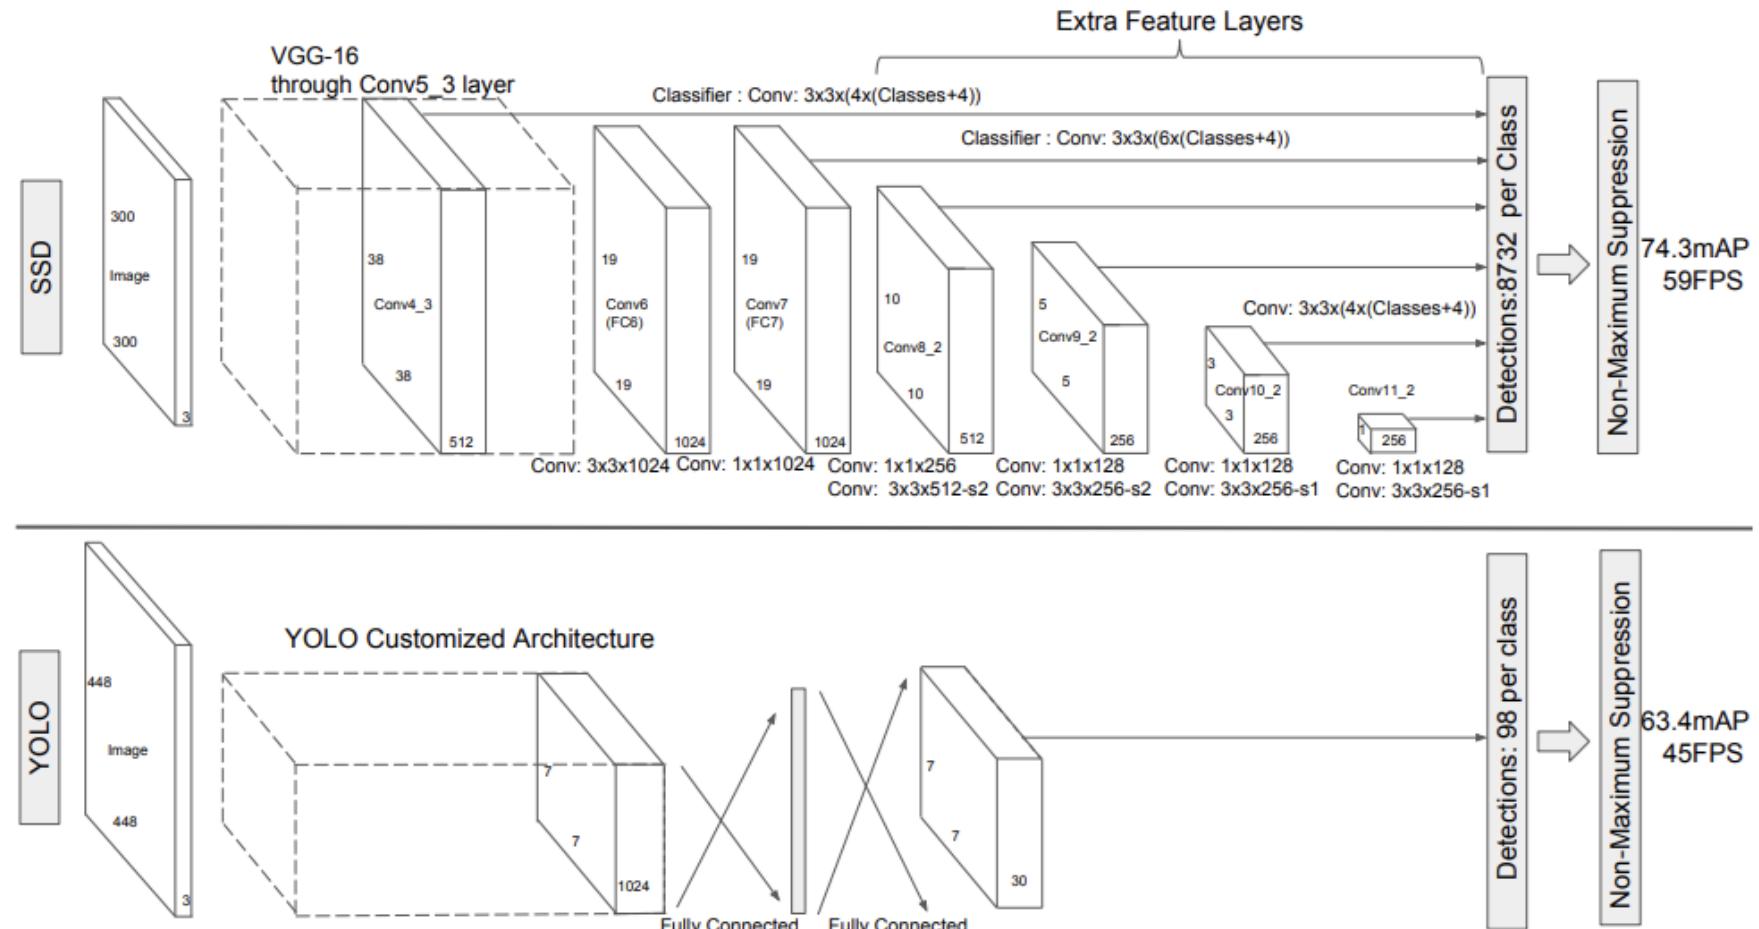


Fig. 2: A comparison between two single shot detection models: SSD and YOLO [5]. Our SSD model adds several feature layers to the end of a base network, which predict the offsets to default boxes of different scales and aspect ratios and their associated confidences. SSD with a  $300 \times 300$  input size significantly outperforms its  $448 \times 448$  YOLO counterpart in accuracy on VOC2007 test while also improving the speed.

- One-stage detector: open-sourced repos

SSD Demo

[https://github.com/hli2020/object\\_detection#testing-ssd](https://github.com/hli2020/object_detection#testing-ssd)

Or

(ipython notebook例子)

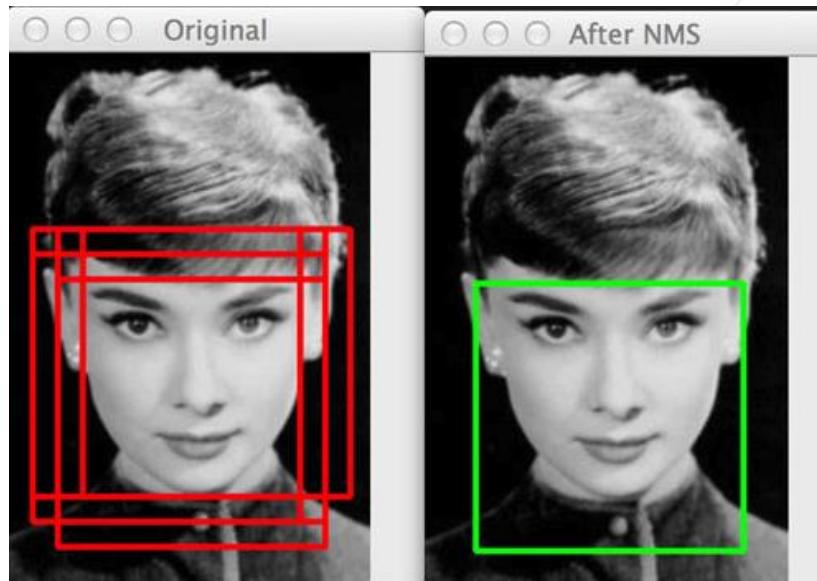
<https://github.com/amdegroot/ssd.pytorch/blob/master/demo/demo.ipynb>

How to implement a YOLO (v3) object detector from scratch in PyTorch

<https://blog.paperspace.com/how-to-implement-a-yolo-object-detector-in-pytorch/>

SSD, YOLO这些方法都是one-stage detector.  
没有RPN过程，直接生成检测结果。

- NMS



一种post-processing 方式。  
用在**所有**检测系统里。

物体检测的指标里，不允许出现  
多个重复的检测，即使这些结果  
和真值都比较近。

那么如何删除多余的检测结果呢？  
**Non-maximum suppression (NMS)**

做法：

把所有检测结果按照分值(conf. score)从高到底排序，保留最高分数的box  
，那么和它距离上最近的那个box, 就没有必要保留了。

以此类推。

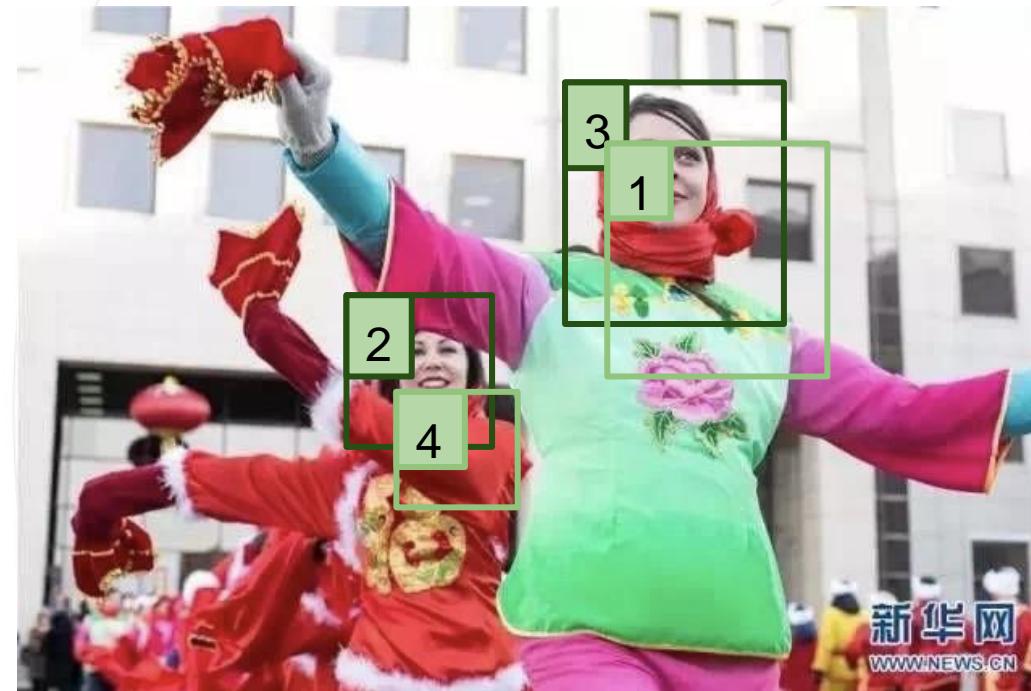
- NMS

按照类别来做的。

右图例子（检测人脸），

1-4分别是分数由高到低的4个目标框，假设1，3被判为距离较近，  
2，4距离很近，

**哪些框保留，哪些要删除？**



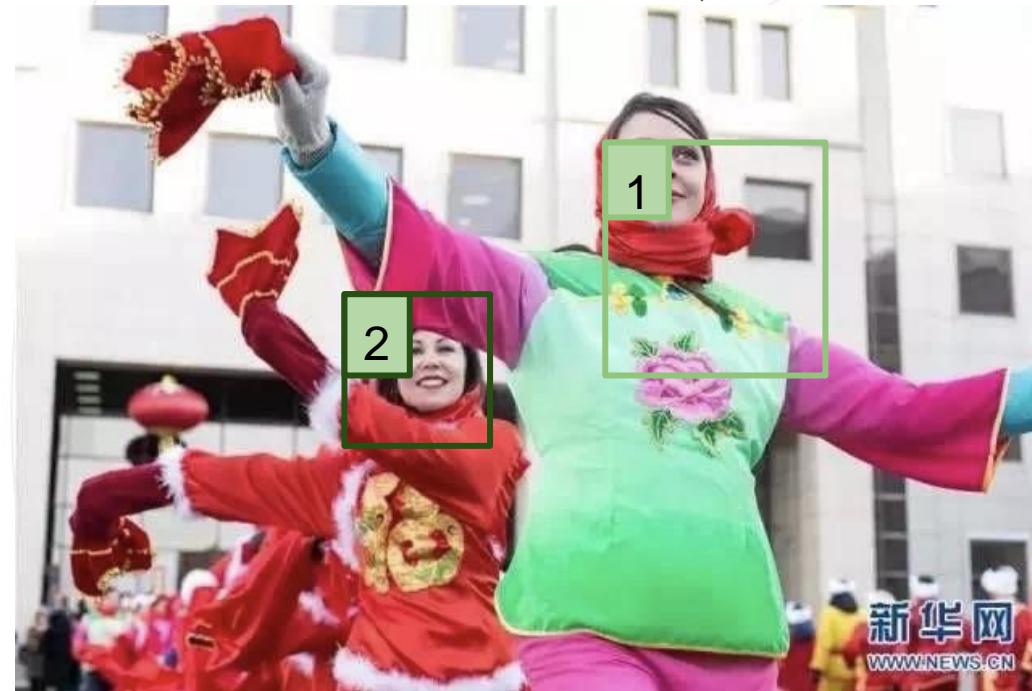
# Recap: Image Detection

- NMS

NMS是按照每一个类别做的

[https://github.com/hli2020/feature\\_intertwiner/blob/master/lib/layers.py#L664](https://github.com/hli2020/feature_intertwiner/blob/master/lib/layers.py#L664)

3, 4 removed!



# Outline

---

**Part 1**      **Recap: classification loss and detection pipeline**

---

**Part 2**      **3D Detection and BEV Perception**

---

**Part 3**      **Image segmentation**

---

# Bird's-eye-view Perception



清华大学  
Tsinghua University

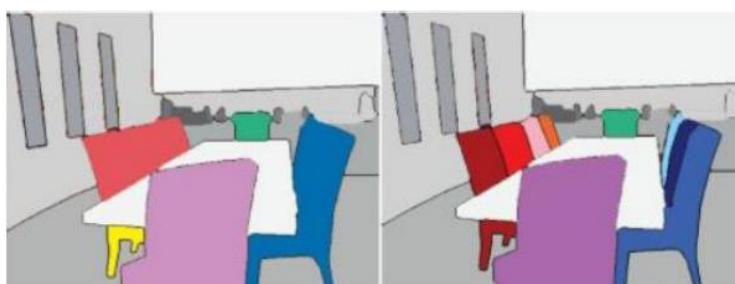
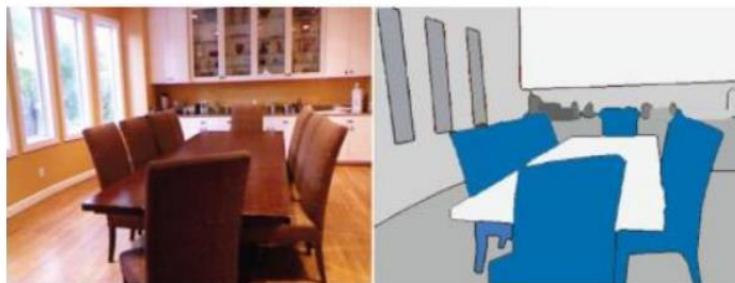


# Outline

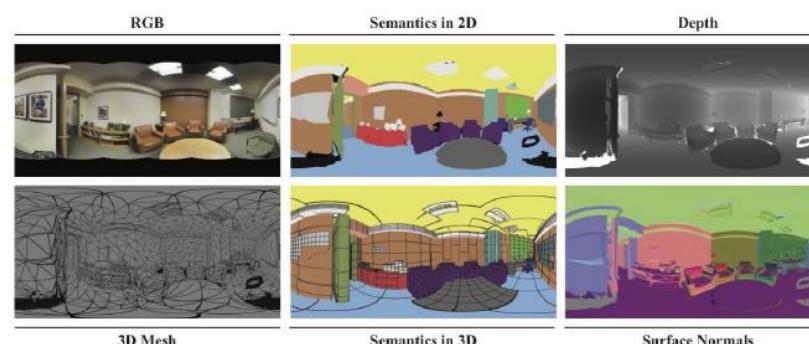
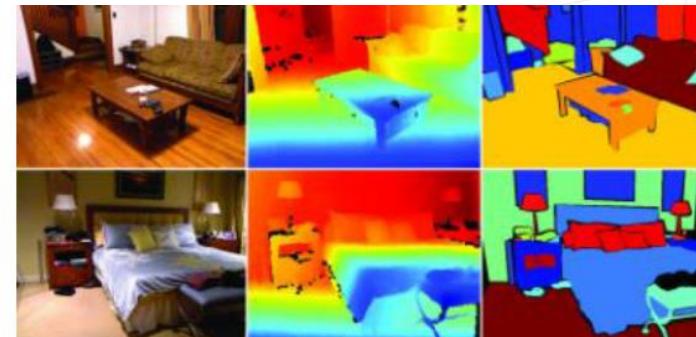
- 
- Part 1**      **Recap: classification loss and detection pipeline**
  - Part 2**      **3D Detection and BEV Perception**
  - Part 3**      **Image segmentation**
-

# Image Segmentation: An Introduction

语义(semantic) or 实例



2D, 2.5D (depth est.) and 3D



3D Mesh

Semantics in 3D

Surface Normals

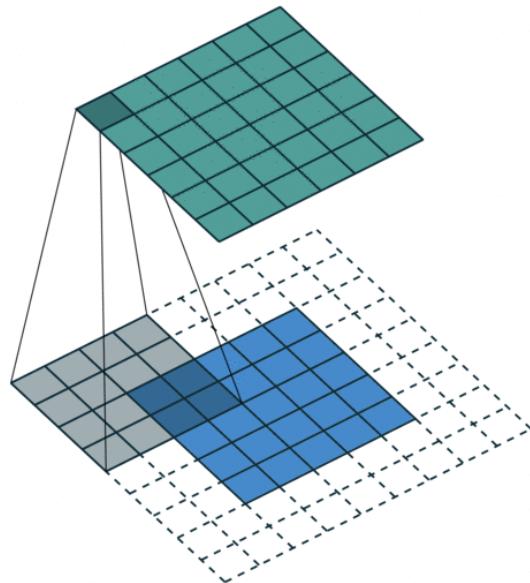
## Semantic segmentation

- FCN
- SegNet
- Dilation
- DeconvNet
- ENet (速度快)
- Deeplab V1 V2 V3
- ParseNet
- RefineNet
- Large Kernel Matters

## Instance segmentation

- SDS
- DeepMask
- SharpMask
- MultiPathNet
- MNC
- Mask-RCNN

## 卷积 (Convolution)



`Class torch.nn.Conv2d(in_channels, out_channels, kernel_size,  
stride=1, padding=0, dilation=1, groups=1, bias=True)`

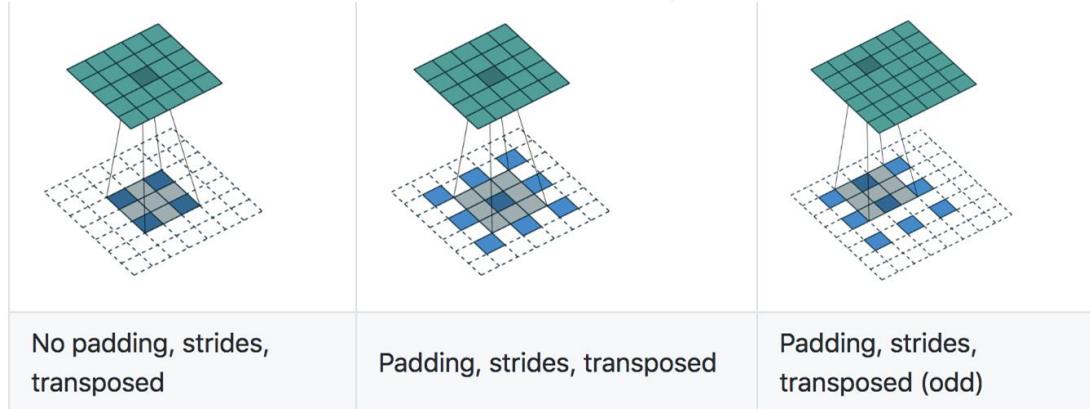
<https://pytorch.org/docs/stable/nn.html#torch.nn.Conv2d>

输出大小是多少？

Output size = 20, 33, 25, 50

公式:  $(W - \text{kernel} + 2 \text{ pad}) / \text{stride} + 1$  向下取整

## 反卷积 (Deconvolution) - upsample



### 普通卷积

$$W_{out} = (W - \text{kernel} + 2 \text{ pad}) / \text{stride} + 1$$

### 反函数

$$W = (W_{out} - 1) * \text{stride} - 2\text{pad} + \text{kernel}$$

### 反卷积公式

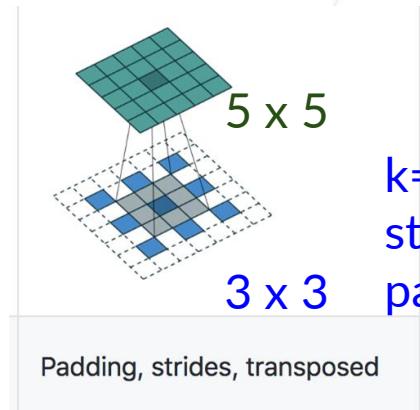
$$W_{out} = (W_{in} - 1) * \text{stride} - 2\text{pad} + \text{kernel}$$

蓝色的是输入 feature map (较小), 绿色的是输出 (较大)

有 stride 版本的反卷积是  
先 up-sample 输入(蓝色), 然后移动 filter, 正常卷积, 得出结果

## 反卷积 (Deconvolution) - upsample

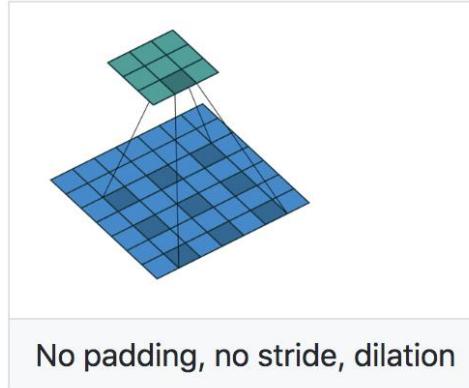
k=3, 问padding 和  
stride是多少?



$$W_{out} = (W_{in} - 1) * \text{stride} - 2\text{pad} + \text{kernel}$$

*Class* `torch.nn.ConvTranspose2d (in_channels,  
out_channels, kernel_size, stride=1, padding=0,  
output_padding=0, groups=1, bias=True, dilation=1)`

## 空洞卷积 (Dilated convolution) - 正常卷积(downsample)的一个细节



- Input:  $(N, C_{in}, H_{in}, W_{in})$
- Output:  $(N, C_{out}, H_{out}, W_{out})$  where

默认值1, 即没有dilation

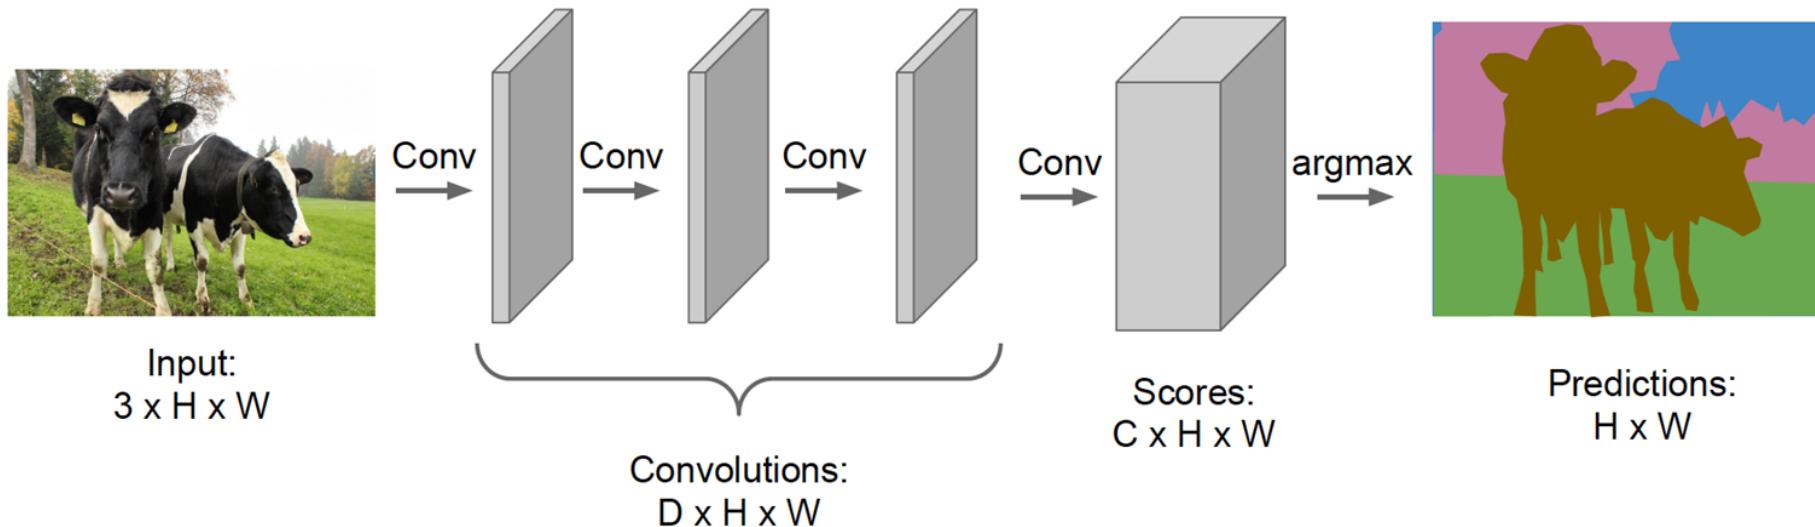
$$H_{out} = \left\lfloor \frac{H_{in} + 2 \times \text{padding}[0] - \text{dilation}[0] \times (\text{kernel\_size}[0] - 1) - 1}{\text{stride}[0]} + 1 \right\rfloor$$
$$W_{out} = \left\lfloor \frac{W_{in} + 2 \times \text{padding}[1] - \text{dilation}[1] \times (\text{kernel\_size}[1] - 1) - 1}{\text{stride}[1]} + 1 \right\rfloor$$

## 稀疏化filter - 扩大视野(receptive field)

注意：和反卷积不同！

# Semantic Segmentation: a naïve approach

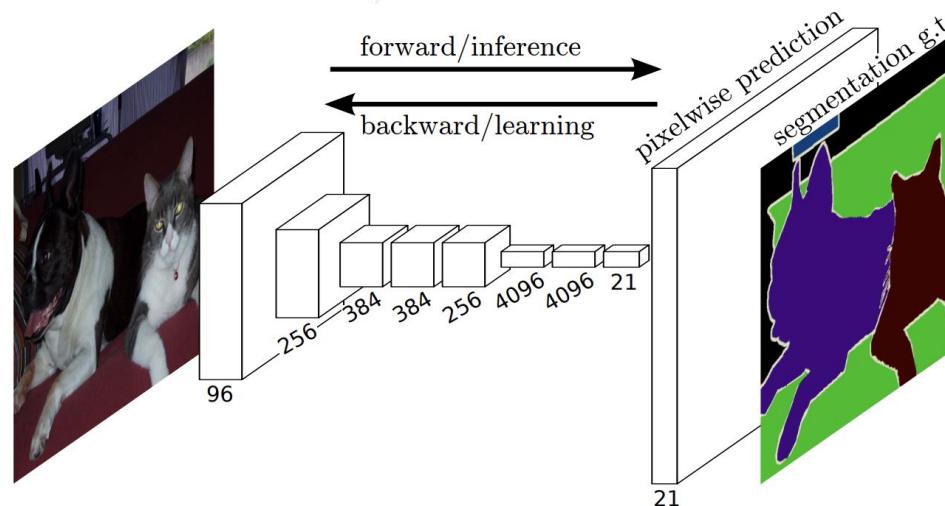
应用一堆卷积，保持feature map和原始图片一致。(channel=D)  
最后一层卷积的channel个数是类别个数即可。(channel=C)



但很显然，计算量较大

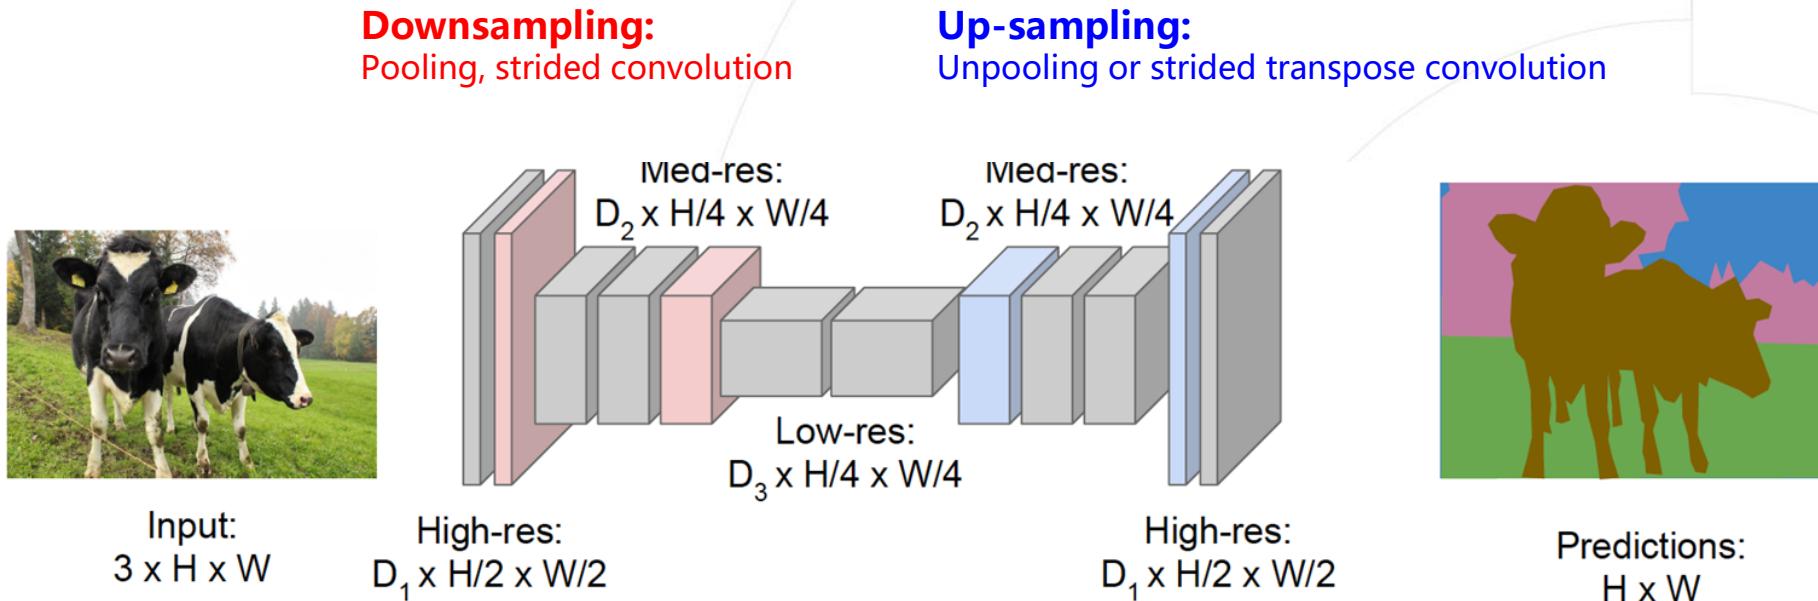
This work is quite like the milestone of RCNN in detection

1. 训练问题：端到端学习
2. 连接层问题：全连接改为全卷积，支持可变输入
3. 特征图变小问题：利用反卷积向上放大特征图
4. 特征融合问题：利用skip connection融合多层特征提高上采样精细度



Long, Shelhamer, and Darrell, “**Fully Convolutional Networks** for Semantic Segmentation”, CVPR 2015  
Noh et al, “Learning Deconvolution Network for Semantic Segmentation”, ICCV 2015

This work is quite like the milestone of RCNN in detection



Long, Shelhamer, and Darrell, “**Fully Convolutional Networks** for Semantic Segmentation”, CVPR 2015  
Noh et al, “Learning Deconvolution Network for Semantic Segmentation”, ICCV 2015

## V1: ICLR 2015

<https://arxiv.org/pdf/1412.7062.pdf>

**Semantic Image Segmentation with Deep Convolutional Nets and Fully Connected CRFs**

Code: <https://bitbucket.org/deeplab/deeplab-public/src/master/>

## V2: arXiv:1606.00915

**DeepLab: Semantic Image Segmentation with Deep Convolutional Nets, Atrous Convolution, and Fully Connected CRFs**

Code: <https://bitbucket.org/aquariusjay/deeplab-public-ver2/src/master/>

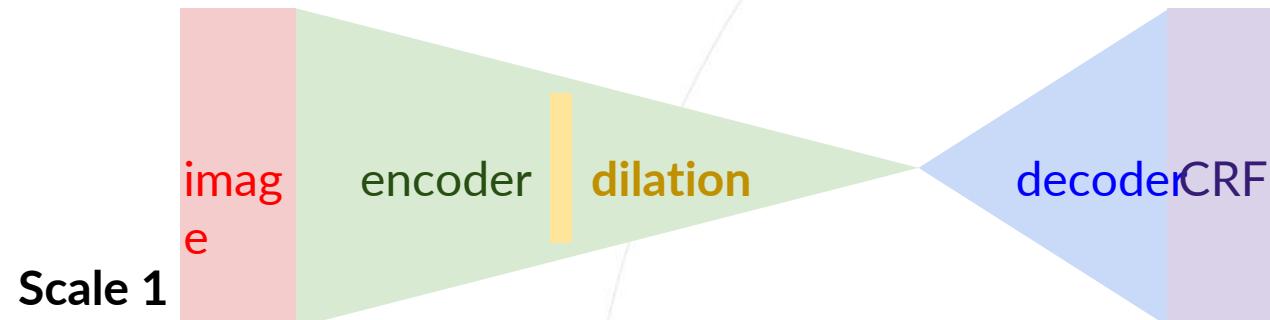
## V3: Rethinking Atrous Convolution for Semantic Image Segmentation

<https://arxiv.org/pdf/1706.05587.pdf>

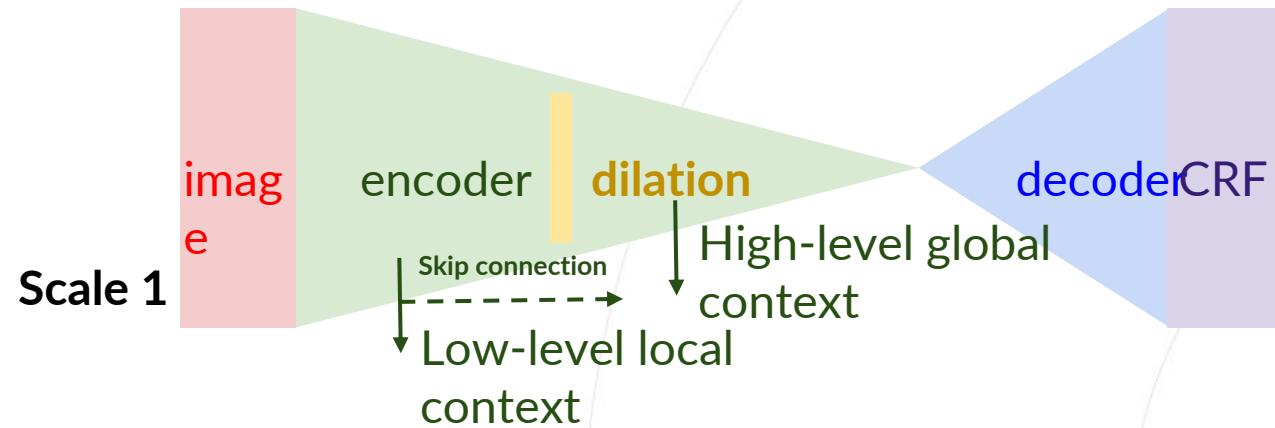
Website: <http://liangchihchen.com/projects/DeepLab.html>

Blog: <https://towardsdatascience.com/review-deeplabv3-atrous-convolution-semantic-segmentation-6d818bfd1d74>

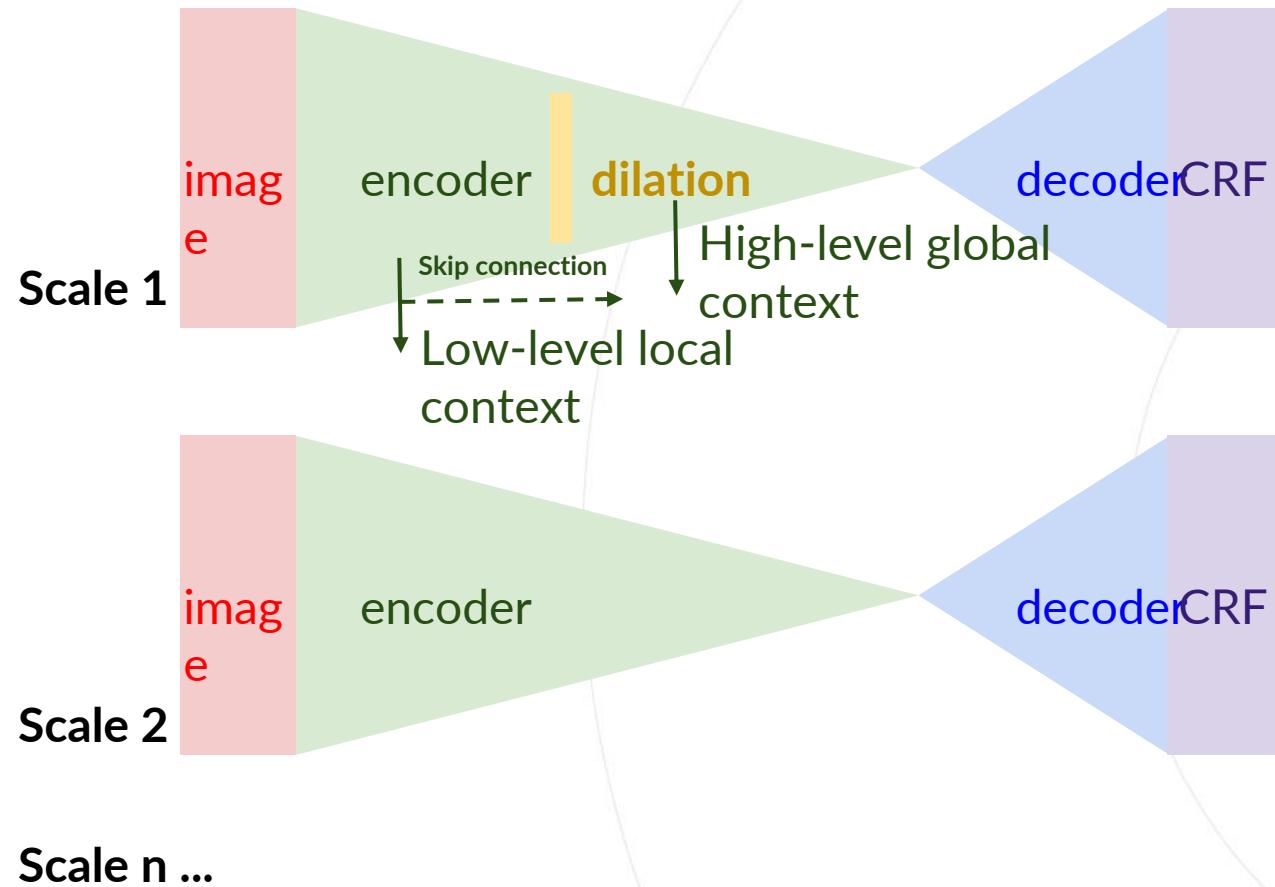
# Semantic Segmentation: general pipeline



# Semantic Segmentation: general pipeline

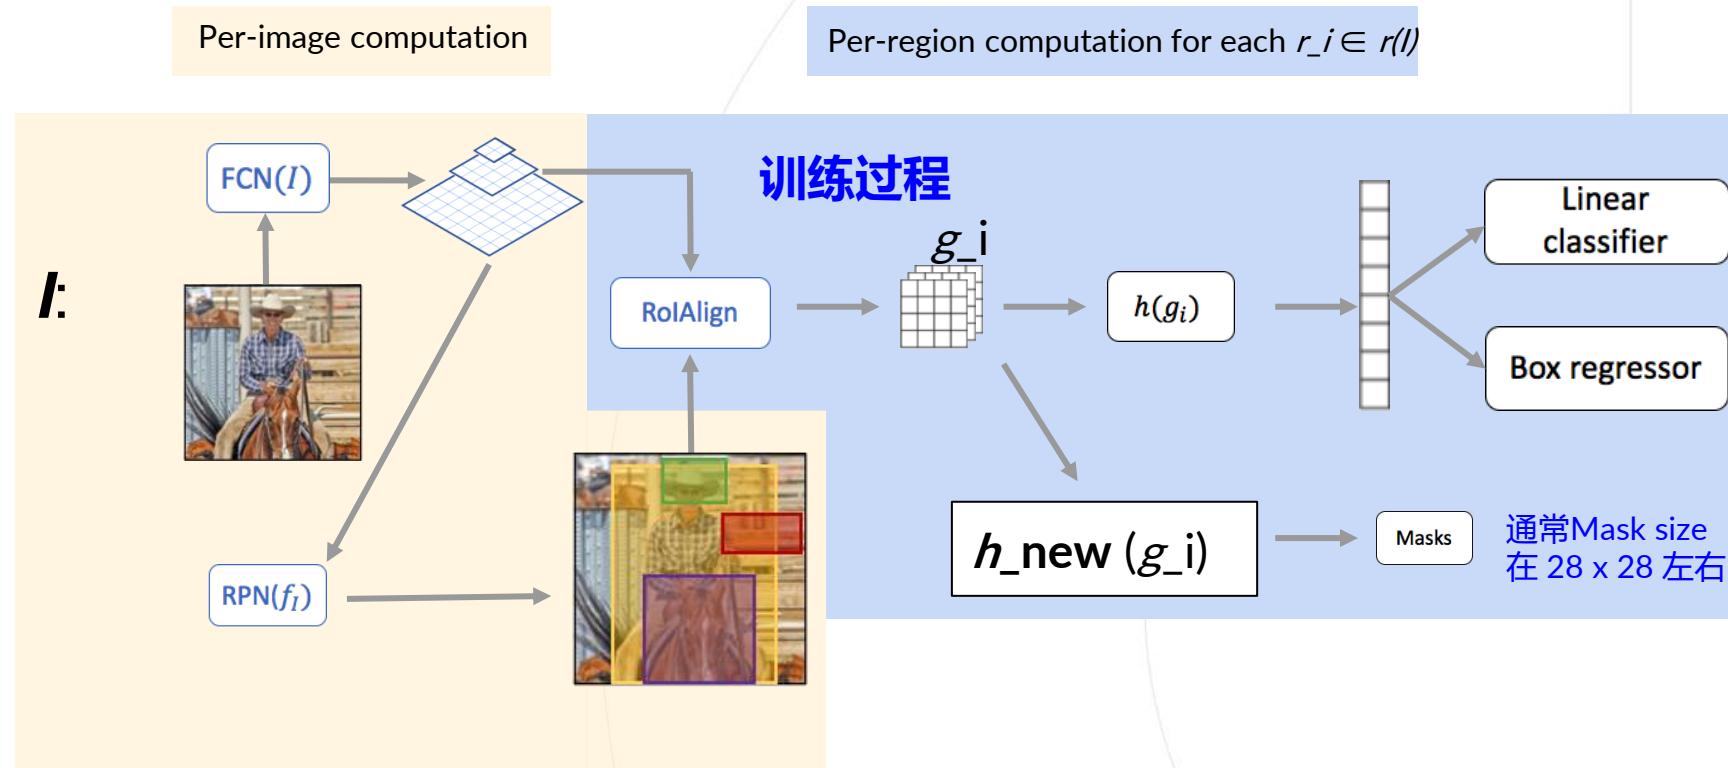


# Semantic Segmentation: general pipeline

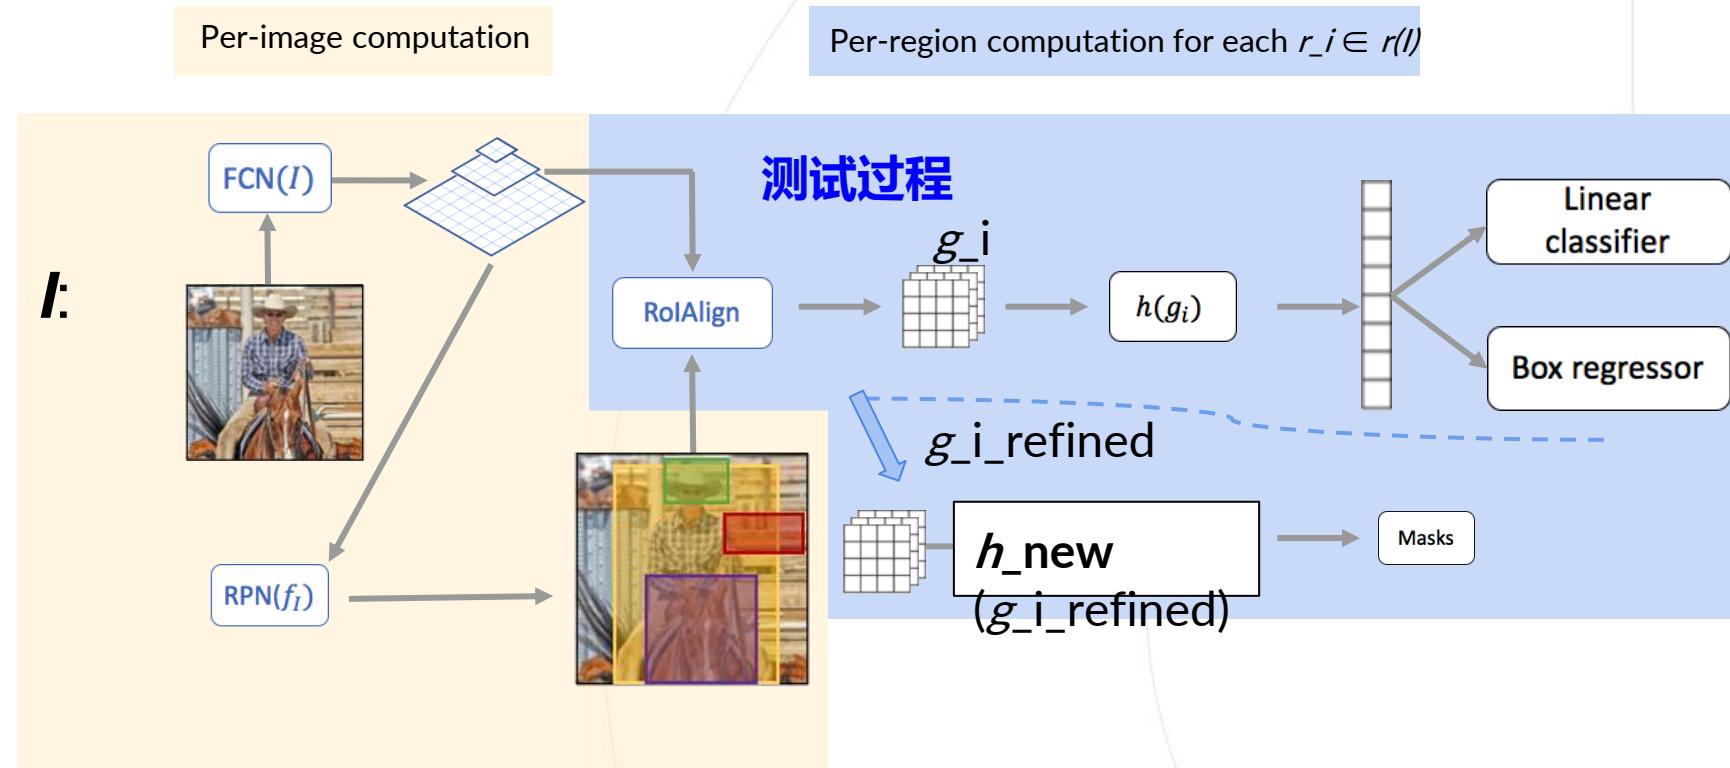


1. 上采样问题      SegNet, DeconvNet, SharpMask, RefineNet
  - a. Encoder-decoder
  - b. Deconvolution
  - c. Unpooling
  - d. Interpolation
2. 底层特征融合      U-net/hourglass structure in pose estimation
3. Receptive field      DeepLab, ParseNet, PSPNet/ICNet
  - a. Dilation /hole
  - b. Global pooling
4. 多尺度      DeepLab
  - a. Multi-scale train/test
  - b. high/low layer feature fusion
  - c. Spatial pyramid pooling

# Instance Segmentation: Mask RCNN

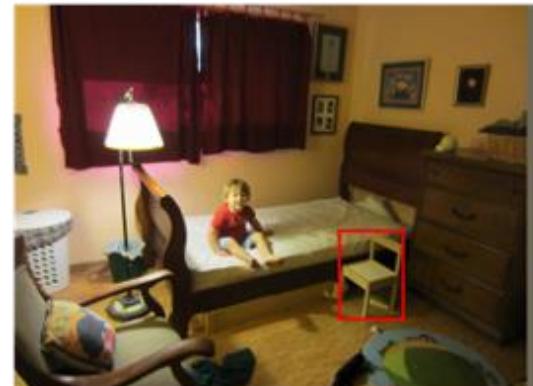


# Instance Segmentation: Mask RCNN



# Instance Segmentation: Mask RCNN - results

Image with training proposal



28x28 mask target

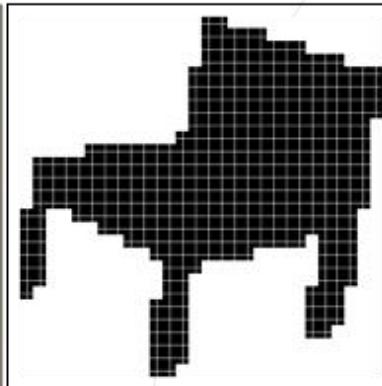
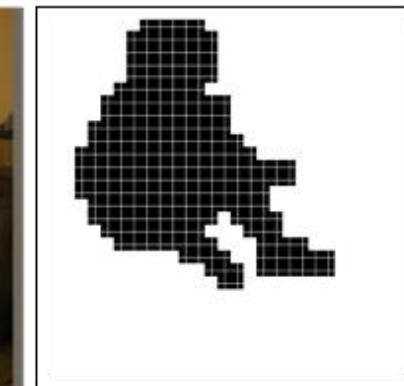
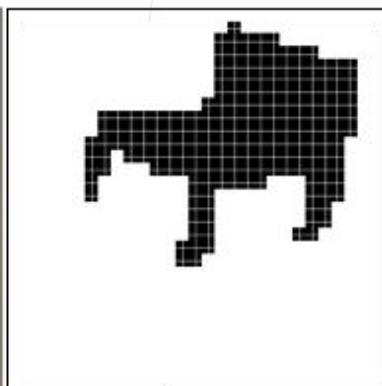
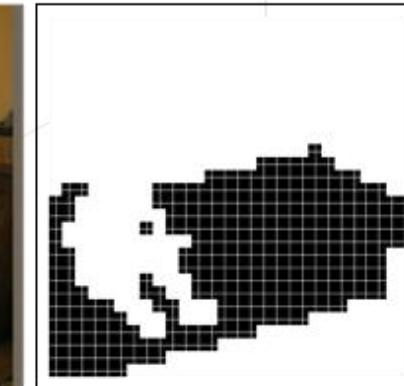


Image with training proposal



28x28 mask target



# Instance Segmentation: Mask RCNN - results



# Image Segmentation: datasets

TABLE 1: Popular large-scale segmentation datasets.

Name and Reference	Purpose	Year	Classes	Data	Resolution	Sequence	Synthetic/Real	Samples (training)	Samples (validation)	Samples (test)
PASCAL VOC 2012 Segmentation [27]	Generic	2012	21	2D	Variable	✗	R	1464	1449	Private
PASCAL-Context [28]	Generic	2014	540 (59)	2D	Variable	✗	R	10103	N/A	9637
PASCAL-Part [29]	Generic-Part	2014	20	2D	Variable	✗	R	10103	N/A	9637
SBD [30]	Generic	2011	21	2D	Variable	✗	R	8498	2857	N/A
Microsoft COCO [31]	Generic	2014	+80	2D	Variable	✗	R	82783	40504	81434
SYNTHIA [32]	Urban (Driving)	2016	11	2D	960 × 720	✗	S	13407	N/A	N/A
Cityscapes (fine) [33]	Urban	2015	30 (8)	2D	2048 × 1024	✓	R	2975	500	1525
Cityscapes (coarse) [33]	Urban	2015	30 (8)	2D	2048 × 1024	✓	R	22973	500	N/A
CamVid [34]	Urban (Driving)	2009	32	2D	960 × 720	✓	R	701	N/A	N/A
CamVid-Sturgess [35]	Urban (Driving)	2009	11	2D	960 × 720	✓	R	367	100	233
KITTI-Layout [36] [37]	Urban/Driving	2012	3	2D	Variable	✗	R	323	N/A	N/A
KITTI-Ros [38]	Urban/Driving	2015	11	2D	Variable	✗	R	170	N/A	46
KITTI-Zhang [39]	Urban/Driving	2015	10	2D/3D	1226 × 370	✗	R	140	N/A	112
Stanford background [40]	Outdoor	2009	8	2D	320 × 240	✗	R	725	N/A	N/A
SiftFlow [41]	Outdoor	2011	33	2D	256 × 256	✗	R	2688	N/A	N/A
Youtube-Objects-Jain [42]	Objects	2014	10	2D	480 × 360	✓	R	10167	N/A	N/A
Adobe's Portrait Segmentation [26]	Portrait	2016	2	2D	600 × 800	✗	R	1500	300	N/A
MINC [43]	Materials	2015	23	2D	Variable	✗	R	7061	2500	5000
DAVIS [44] [45]	Generic	2016	4	2D	480p	✓	R	4219	2023	2180
NYUDv2 [46]	Indoor	2012	40	2.5D	480 × 640	✗	R	795	654	N/A
SUN3D [47]	Indoor	2013	—	2.5D	640 × 480	✓	R	19640	N/A	N/A
SUNRGBD [48]	Indoor	2015	37	2.5D	Variable	✗	R	2666	2619	5050
RGB-D Object Dataset [49]	Household objects	2011	51	2.5D	640 × 480	✓	R	207920	N/A	N/A
ShapeNet Part [50]	Object/Part	2016	16/50	3D	N/A	✗	S	31,963	N/A	N/A
Stanford 2D-3D-S [51]	Indoor	2017	13	2D/2.5D/3D	1080 × 1080	✓	R	70469	N/A	N/A
3D Mesh [52]	Object/Part	2009	19	3D	N/A	✗	S	380	N/A	N/A
Sydney Urban Objects Dataset [53]	Urban (Objects)	2013	26	3D	N/A	✗	R	41	N/A	N/A
Large-Scale Point Cloud Classification Benchmark [54]	Urban/Nature	2016	8	3D	N/A	✗	R	15	N/A	15

Pascal, COCO, Cityspace (cars and all),  
KITTI

假设 $K+1$ 类，下标从0到 $k$ 。 $p_{ij}$  表示属于 $i$ 类的样本被预测为 $j$ 类。

Pixel Accuracy(PA) = (预测对的像素个数)/(总的像素个数)

Mean Pixel Accuracy(MPA)= 平均每类的准确率

Mean IoU=平均每类的IOU

$$MIoU = \frac{1}{k+1} \sum_{i=0}^k \frac{p_{ii}}{\sum_{j=0}^k p_{ij} + \sum_{j=0}^k p_{ji} - p_{ii}}$$

Frequency weighted IoU=加权后每类的IOU

$$PA = \frac{\sum_{i=0}^k p_{ii}}{\sum_{i=0}^k \sum_{j=0}^k p_{ij}}$$

$$MPA = \frac{1}{k+1} \sum_{i=0}^k \frac{p_{ii}}{\sum_{j=0}^k p_{ij}}$$

$$FWIoU = \frac{1}{\sum_{i=0}^k \sum_{j=0}^k p_{ij}} \sum_{i=0}^k \frac{\sum_{j=0}^k p_{ij} p_{ii}}{\sum_{j=0}^k p_{ij} + \sum_{j=0}^k p_{ji} - p_{ii}}$$

# LiDAR Semantic Segmentation



清华大学  
Tsinghua University





清华大学  
Tsinghua University



# END

Reach me at [lihongyang@senseauto.com](mailto:lihongyang@senseauto.com)